

AATIF: A Multiplicative Governance Equation for Arabic-First LLM Safety

Abdulmjeed Ibrahim Khenkar
Independent Researcher
abdulmjeed.ibrahem@gmail.com

June 2026

Abstract

We present AATIF, a mathematical governance framework for large language model safety that provides an explicit, interpretable safety equation with provable non-compensability. The core governance function $S = \sigma(w_1 \cdot I + w_2 \cdot E) \cdot [1 - \sigma(\alpha(H - \theta))]$ fuses three independent perception channels—semantic harm proximity (H), intent classification (I), and emotional load (E)—through a multiplicative gate structure that prevents benign intent from compensating for detected harm. Unlike trained classifiers, AATIF requires zero training data, operating through 169 curated semantic anchors for harm detection with cosine similarity scoring via multilingual embeddings (bge-m3). The framework treats Arabic as a first-class semantic input, with 35 dialect-hyperbole disambiguation anchors to prevent figurative speech from triggering false harm signals. Evaluated against HarmBench (236 behaviors) and MultiJail (75 prompts), AATIF achieves 74.3% safety-category detection with 100% on chemical/biological threats, and Arabic outperforms English on identical harmful prompts (74.7% vs 69.3%), validating the Arabic-first design. The system includes hysteresis control for decision stability, unconditional CBRN safety gates, and four tunable mode profiles calibrated via A/B testing (in-sample $F_1 = 0.984$ on 54 cases; held-out $F_1 = 0.9615$ on 56 unseen cases after anchor refinement). A test suite of 166 deterministic tests ensures mathematical correctness. AATIF combines, for the first time in a single deployable system, a closed-form multiplicative governance equation with mathematically guaranteed non-compensability, Arabic dialect-aware anchor scoring, and zero-fine-tuning extensibility, designed for human-over-the-loop governance.

Keywords: LLM safety, governance equation, Arabic NLP, multiplicative gating, non-compensability, semantic anchors

1 Introduction

Large language model (LLM) safety systems share three structural limitations that persist across both commercial platforms and academic proposals.

Gap 1: No interpretable equation. No deployed safety system provides an explicit, interpretable mathematical governance formula with provable properties. FanarGuard [Fatehkia et al., 2025] is a 2B-parameter classifier trained on 468K prompt-response pairs; Llama Guard 3 [Inan et al., 2024] is a fine-tuned language model performing binary classification; Google’s Perspective API [Google Jigsaw, 2017] returns per-attribute toxicity scores from a neural classifier. All three produce scores, but none exposes a closed-form equation that a reviewer can inspect, verify, or prove properties about. The gap is not one of performance but of *interpretability*: when a safety system refuses a request, why?

Gap 2: No non-compensability guarantee. Existing aggregation methods—typically additive or mean-based—allow benign signals in one channel to compensate for harm detected in another. A prompt expressing harmful intent in calm, polite language can score as “moderate risk” under additive aggregation, because the low emotional-load score partially offsets the high harm score. This *toxic positivity* attack vector is a structural property of additive scoring, not an implementation bug.

Gap 3: Arabic as afterthought. Arabic dialect awareness is either absent or retrofitted onto English-first architectures. Arabic colloquial speech routinely uses violent metaphors for everyday emotions—*wallāhi ba-dhbāḥuh* (“I swear I’ll slaughter him”) means *I’m annoyed at him*—and English-first embedding spaces collapse these figurative expressions into their literal semantic neighborhoods, producing systematic false positives. No existing system treats Arabic dialectal hyperbole as a first-class safety concern.

This paper presents AATIF (Architected Adaptive Thoughts & Intelligence Frameworks), a governance framework that addresses all three gaps with a single mathematical structure. The core contribution is the *S equation*:

$$S = \sigma(w_1 \cdot I + w_2 \cdot E) \cdot [1 - \sigma(\alpha(H - \theta))] \quad (1)$$

where H is semantic harm proximity, I is intent classification, E is emotional load, σ is the logistic sigmoid, and θ is the harm threshold. The multiplicative structure ensures *non-compensability*: when H exceeds θ , the harm gate drives S toward zero regardless of how benign the intent or calm the emotional state. A hard override ($H > 0.7 \Rightarrow \text{SAFE_FREEZE}$) provides an absolute safety floor independent of parameter calibration.

AATIF originated from 73 governance conjectures derived through a year of sustained observation across four LLM platforms (ChatGPT, Claude, Gemini, DeepSeek), documented in 78 field notes available as supplementary material. This paper operationalizes one of those conjectures—Conjecture #063, that governance operates through probabilistic constraint density rather than rule retrieval—into a computational engine with an explicit mathematical formulation, a comprehensive test suite of 166 deterministic tests, and benchmark evaluation against HarmBench and MultiJail.

The remainder of this paper is structured as follows. Section 2 reviews related work. Section 3 presents the S equation, its three scorers, the hysteresis controller, safety gates, and mode profiles. Section 4 describes the Arabic-first design. Section 5 reports experimental evaluation. Section 6 discusses contributions and limitations. Section 7 concludes.

2 Related Work

2.1 Alignment Approaches

The dominant paradigm in LLM alignment treats governance as top-down principle application. Christiano et al. [2017] established reinforcement learning from human feedback (RLHF), learning reward functions from pairwise preference comparisons. Ouyang et al. [2022] applied this at scale in InstructGPT, demonstrating that preference-based fine-tuning shifts model outputs toward human-annotated intent. Bai et al. [2022] proposed Constitutional AI, in which natural-language principles guide self-critique and revision. In all three approaches, governance principles precede observation: the direction of derivation runs from principle to model.

A parallel body of work has examined how alignment operates internally. Arditi et al. [2024] demonstrated that refusal behavior across thirteen chat models is mediated by a single linear direction in activation space, rather than by distributed rule-following mechanisms. Lee et al. [2024] showed that direct preference optimization (DPO) does not eliminate harmful

capabilities but learns to route around them, with the original capacity remaining recoverable. Wolf et al. [2024] formalized this fragility through Behavior Expectation Bounds, proving that any alignment process that attenuates rather than removes a behavior remains susceptible to adversarial elicitation. Taken together, these findings suggest that alignment functions as probabilistic constraint density on output distributions rather than as explicit rule retrieval—a characterization that motivates AATIF’s mathematical approach to governance.

2.2 Arabic-Aware AI Safety

FanarGuard [Fatehikia et al., 2025] is the most directly relevant system: a bilingual Arabic-English content moderation classifier developed by the Qatar Center for Artificial Intelligence, trained on 468K prompt-response pairs and fine-tuned into a 2B-parameter model achieving $F_1 = 0.82$ on Arabic safety classification. FanarGuard introduces a dual-axis evaluation—safety (harmful→harmless) and cultural alignment (aligned→unaligned)—representing an important advance in culturally-aware moderation. Llama Guard 3 [Inan et al., 2024] extends Meta’s safety classifier to eight languages including Arabic, applying the MLCommons 13-category taxonomy through binary classification. Google’s Perspective API [Google Jigsaw, 2017] provides per-attribute toxicity scores with Arabic language support. Each of these systems operates as a trained classifier or fine-tuned language model; none provides an explicit, interpretable governance equation.

In Arabic hate speech detection, ADHAR [Aldjanabi et al., 2024] is a multi-dialectal corpus covering five Arabic variants (Egyptian, Levantine, Gulf, Maghrebi, MSA) with inter-annotator agreement $\kappa = 0.95$; cross-dialect experiments show accuracy dropping from 0.84 to 0.64 when training on one dialect and testing on another, empirically demonstrating that dialect-awareness is not optional for Arabic safety systems. SOD [Asiri & Saleh, 2024] contributes a Saudi-specific offensive language corpus of 24K+ tweets with hierarchical annotation, achieving $F_1 = 0.87$ with AraBERT. Both datasets informed AATIF’s anchor expansion for dialectal coverage.

AATIF differs from these systems in three structural respects: (1) it provides an explicit mathematical equation with provable properties rather than learned classifier weights, (2) its multiplicative gate prevents intent from compensating for harm—a non-compensability guarantee that additive aggregation cannot provide, and (3) it requires zero training data, operating through curated semantic anchors. These differences position AATIF as complementary to data-driven classifiers rather than competitive with them.

2.3 Safety Benchmarks

Two benchmark suites are relevant to evaluating AATIF. HarmBench [Mazeika et al., 2024] provides 236 harmful behaviors across seven semantic categories (chemical/biological, cyber-crime, illegal activities, harassment, harmful content, misinformation, copyright). HarmBench is English-only, creating a language mismatch with AATIF’s Arabic-first anchor design that tests cross-lingual transfer. MultiJail [Deng et al., 2024] addresses the multilingual gap with 75 harmful prompts translated by native speakers into 10 languages including Arabic, enabling direct comparison of detection performance across languages. The MLCommons AI Safety taxonomy [MLCommons, 2024] provides a standardized 13-category harm classification (S1–S13) serving as a common reference for coverage comparison.

3 The AATIF Governance Equation

3.1 Definition and Properties

The AATIF governance engine computes a composite safety score $S \in [0, 1]$ from three independent perceptual channels:

$$S = \sigma(w_1 \cdot I + w_2 \cdot E) \cdot \underbrace{[1 - \sigma(\alpha(H - \theta))]}_{\text{harm gate}} \quad (2)$$

The equation consists of two factors whose product determines the governance decision.

Quality term. $\sigma(w_1 \cdot I + w_2 \cdot E) \in (0, 1)$ assesses how clear and benign the request is. Intent classification I captures whether the request exhibits harmful planning signals (higher I = more benign); emotional load E captures whether the user’s language exhibits distress or escalation markers (higher E = calmer state). Parameters $w_1 = 2.0$ and $w_2 = 1.5$ weight intent more heavily than emotion, reflecting the design judgment that *what* the user wants matters more than *how* they feel when asking.

Harm gate. $[1 - \sigma(\alpha(H - \theta))] \in (0, 1)$ is a soft gate controlled by semantic harm proximity H . The steepness parameter $\alpha = 10$ controls transition sharpness; the threshold $\theta = 0.40$ sets the harm level at which the gate begins to close. When $H \ll \theta$, the gate is fully open (≈ 1); when $H \gg \theta$, the gate closes (≈ 0), driving S toward zero.

Non-compensability. The *product* of these two factors is the critical design choice. Because S is the product of two terms each bounded in $(0, 1)$, high harm closes the gate regardless of how benign the intent or calm the emotional state. There exists no combination of I and E values that can force S above the EXECUTE threshold when $H > \theta$. Additive aggregation ($S = aH + bI + cE$) lacks this property: sufficiently high I and E can always compensate for moderate H .

Sigmoid symmetry. The harm gate satisfies $[1 - \sigma(\alpha(H - \theta))] = \sigma(\alpha(\theta - H))$, verified formally. This symmetry ensures the gate is centered exactly at θ and monotonically decreasing in H .

Hard override. When $H > 0.7$, the system triggers `SAFE_FREEZE` regardless of S , providing an absolute safety floor independent of parameter calibration. This override is not derived from the equation but imposed externally, ensuring that no parameter tuning can disable the safety floor.

Decision mapping. The continuous score S maps to four discrete actions:

Table 1: Decision mapping from S score to governance action.

Condition	Action	Behavior
$S > 0.7$	EXECUTE	Process request normally
$0.5 < S \leq 0.7$	CLARIFY	Request clarification before proceeding
$0.3 < S \leq 0.5$	SAFE_STOP	Decline with explanation
$S \leq 0.3$	SAFE_FREEZE	Maximum caution (score-based)
$H > 0.7$	SAFE_FREEZE	Unconditional refusal (hard override)

Follow-up signal. A follow-up signal F provides a harm-floor guarantee: even when S is high, F flags residual risk for downstream monitoring, ensuring that borderline cases receive continued attention.

3.2 The Three Scorers

Each scorer operates independently, producing a continuous score in $[0, 1]$ through cosine similarity against curated semantic anchors using multilingual embeddings (bge-m3 via Ollama). A TF-IDF character n-gram fallback is available for offline operation without GPU.

H scorer—Semantic Harm Proximity (*harārat al-kalima*). The H scorer maintains 169 curated anchors spanning violence, weapons, self-harm, cybercrime, hate speech, and chemical/biological categories. Each anchor is a short phrase representing a harm concept. Input text is embedded via bge-m3 and compared against all anchors via cosine similarity. *Top-K aggregation* ($K = 3$) averages the three highest-similarity scores, preventing any single anchor from dominating the result. This mitigates the *anchor contamination* effect, in which one high-similarity anchor inflates the score despite low agreement across the anchor set. The H scorer includes 35 dialect-specific hyperbole anchors scored at $H \leq 0.10$, providing semantic “counter-gravity” for figurative Arabic expressions (Section 4).

I scorer—Intent Classification (*al-niyya*). The I scorer maintains 44 intent anchors organized into categories: informational, transactional, navigational, creative, and harmful-planning. Input is classified via top- K softmax over anchor similarities, with an out-of-distribution (OOD) guard that flags inputs dissimilar to all intent categories. Confidence bands map similarity scores to discrete confidence levels.

E scorer—Emotional Load (*al-shu‘ūr*). The E scorer maintains 32 emotion anchors covering distress, escalation, calmness, and neutral states. Higher E indicates calmer emotional state; distress and escalation indicators produce lower E scores, which reduce S and may shift the decision toward CLARIFY. The E scorer partially implements a *textual spectrogram*—a reconstruction of emotional information from textual features when prosodic cues are absent (Section 4.3).

All three scorers are deterministic at test time: given the same input and anchor set, they produce identical scores across runs.

3.3 Hysteresis Controller

When S hovers near a decision boundary, the system can produce unstable decision sequences—oscillating between EXECUTE and CLARIFY on successive evaluations of similar inputs. The hysteresis controller addresses this by introducing entry/exit threshold offsets: transitioning from EXECUTE to CLARIFY requires S to fall below the boundary by offset ε , while returning to EXECUTE requires S to rise above it by the same margin. This is the “AC thermostat” principle: a heating system that turns on at 68°F and off at 72°F is more stable than one that toggles at exactly 70°F. The controller maintains per-session state, ensuring independent hysteresis tracking across concurrent user conversations.

3.4 Safety Gates: Law Ω and Law Ξ

Two unconditional safety gates fire *before* the S equation is evaluated, providing defense-in-depth independent of the scoring mechanism.

Law Ω (CBRN Gate). Unconditional detection of chemical, biological, radiological, and nuclear (CBRN) content in both Arabic and English. When triggered, the system enters SAFE_FREEZE regardless of all other scores. This gate is mode-independent: it fires identically across all four mode profiles.

Law Ξ (Override Lock). Detection of jailbreak and override attempts—phrases such as “ignore all previous instructions” or *tajāwuz al-ta‘līmāt* (“override instructions”). When triggered, the system enters `SAFE_FREEZE`. The override lock is designed to be non-bypassable: no prompt construction can deactivate it from within the input channel.

Both gates operate on pattern matching against curated trigger lists and are independent of the embedding backend, ensuring they function even in TF-IDF fallback mode.

3.5 Mode Profiles

AATIF supports four parameter profiles that adjust the sensitivity of the S equation without changing its structure:

Table 2: Mode profiles with parameter settings and use cases.

Profile	α	θ	Use case
<code>high_sensitivity</code>	15	0.45	High-risk domains (medical, legal)
<code>default</code>	10	0.40	Standard operation
<code>creative</code>	8	0.55	Creative writing, fiction
<code>balanced_strict</code>	10	0.40	Tighter gate (A/B calibrated)

Higher α makes the gate transition sharper (less ambiguous region); higher θ makes the gate more tolerant (opens wider before closing). The `high_sensitivity` profile uses a steeper gate ($\alpha = 15$) and reduced emotion weight ($w_2 = 1.0$), producing a narrower transition zone and less emotional compensation—achieving conservative behavior through gate sharpness rather than threshold placement. The `creative` profile tolerates higher similarity before triggering ($\theta = 0.55$), with a gentler transition ($\alpha = 8$). Crucially, the hard override ($H > 0.7$) and safety gates (Laws Ω , Ξ) are mode-independent—they fire at full sensitivity regardless of profile selection.

3.6 Ambiguity Pre-Check and Jailbreak Detection

Ambiguity detection. Short or vague prompts trigger an override from `EXECUTE` to `CLARIFY`, preventing the system from processing ambiguous requests that could be interpreted as either benign or harmful. This operationalizes the design principle that the cheapest correct response to an ambiguous input is a clarifying question, not a speculative completion.

Bidirectional jailbreak handler. The jailbreak handler operates in both directions: it *escalates* decisions when jailbreak markers are detected (role-play framing, instruction-override language, system-prompt extraction attempts), and it can *downgrade* a `SAFE_FREEZE` to `SAFE_STOP` when no jailbreak markers are present but simple harm content is detected. This bidirectional design reflects the principle that simple harm (a genuine safety question from a concerned user) is not manipulation and should receive an empathetic explanation rather than unconditional refusal.

4 Arabic-First Design

4.1 The Meaning Density Hypothesis (*kathāfat al-ma‘nā*)

We hypothesize that *meaning density*—the number of semantic dimensions encodable per token—is the relevant variable for selecting a base representation language for a safety system. Arabic was chosen not from cultural preference but from measurable information-theoretic properties that make it the strongest known candidate among major world languages.

Arabic’s root-pattern morphology provides transparent semantic compression. The triliteral root *k-t-b* (“writing”) generates *kitāb* (book), *kātib* (writer), *maktaba* (library), *maktūb* (written letter)—a single root encoding semantic relationships that English requires unrelated lexemes to express. Classical Arabic lexicography documents over 700 names for the lion, each encoding a distinct behavioral dimension (resting, roaring, hunting, stalking)—not synonyms but distinct semantic coordinates in a high-dimensional space. Figurative density compounds this: *wallāhi ba-dhbaḥuh* (“I swear I’ll slaughter him”) packs intent, emotion, cultural context, and literal meaning into two words.

This density has a concrete consequence for safety systems. English-first embedding models project Arabic text into a semantic space whose geometry was built for English lexical distinctions. This projection systematically collapses the high-dimensional Arabic semantic space: the 700 lion-names collapse to a single cluster near “lion”; figurative and literal violence collapse into the same neighborhood. The result is information loss at precisely the semantic boundaries where safety decisions are most critical.

AATIF’s curated anchors are a workaround for this projection loss: Arabic-language anchors provide reference points in the embedding regions where English-first geometry is coarsest. The principled solution—an Arabic-first embedding model trained to preserve Arabic semantic distinctions—is future work. We note that the meaning density hypothesis is not exclusive to Arabic: other Semitic languages (Hebrew, Amharic) share root-pattern morphology and would exhibit analogous properties [Wolff & Holmes, 2011, Slobin, 1996]. Our MultiJail results (Arabic 74.7% vs English 69.3%) provide circumstantial evidence consistent with this hypothesis, though testing it directly requires building an Arabic-first embedding model.

4.2 Dialect Hyperbole Disambiguation

Arabic colloquial speech uses violent metaphors for everyday emotions with a frequency that poses a unique challenge for safety systems. Table 3 illustrates representative examples across Gulf, Levantine, and Egyptian dialects.

Table 3: Arabic dialect hyperbole: literal meaning vs. actual meaning.

Expression	Literal meaning	Actual meaning
<i>yijīb lahum jalṭa</i>	Give them a stroke	Annoy them greatly
<i>wallāhi ba-dhbaḥuh</i>	I’ll slaughter him	I’m annoyed at him
<i>qunbula ‘ala al-sōshyāl mīdyā</i>	A bomb on social media	A viral hit
<i>bamūt fīk</i>	I’m dying in you	I love you deeply
<i>yaqtulnī min al-daḥk</i>	He’s killing me	He’s hilarious

A safety system without dialect awareness would flag all five expressions as harmful: “slaughter,” “bomb,” “dying,” “killing” all appear in violence-related anchor neighborhoods. The H scorer addresses this through 35 dialect-hyperbole anchors, each scored at $H \leq 0.10$. These anchors act as semantic counter-gravity: when an input matches a hyperbole anchor more closely than a harm anchor, the figurative interpretation dominates and the harm score remains low. This mechanism exploits the top- K aggregation design—if one of the top-3 nearest anchors is a hyperbole anchor, it dilutes the average toward safety.

We call the underlying failure mode *Lexical Anchor Contamination*: English-first embeddings cannot reliably separate Arabic figurative violence from literal violence because the embedding geometry was not built to encode this distinction. The cross-dialect accuracy drop documented by Aldjanabi et al. [2024] (0.84 to 0.64) confirms that dialect-sensitivity is not optional.

The dialect test suite contains 22 test cases covering Gulf, Levantine, and Egyptian dialectal expressions. All 22 pass with zero false positives: no hyperbolic expression triggers harm detection.

4.3 The Maqam-to-Safety Origin Chain

The emotional channel (E scorer) and the multiplicative gate structure originated from an unexpected domain: Arabic music theory. We present this as an intellectual origin chain—a sequence of observations leading from quarter-tone music to a concrete safety mechanism.

The quarter-tone gap. Arabic music uses the maqam system [Touma, 1971, Abu Shumays, 2013], built on quarter-tone intervals that cannot be represented by the Western 12-tone equal-tempered scale. AI music systems trained on Western scales are structurally incapable of representing Arabic musical emotion—not due to insufficient training data, but because the representational vocabulary lacks the required resolution. EEG studies confirm that maqam produces distinct neural responses in human listeners [Yaghmour et al., 2021], establishing that the emotional content is real, not merely cultural convention.

From music to speech. Each Arabic maqam carries temporal-emotional information: Bayati evokes sadness, Rast stability, Hijaz longing. The same principle applies to speech: prosody (pitch contour, rhythm, emphasis) carries emotional information that lexical content alone does not capture.

Text loses prosody. When communication shifts from speech to text, prosodic information is lost. The sentence “fine, do whatever you want” can be angry, sarcastic, resigned, or genuinely permissive. Text is prosody-stripped, and safety systems operating on text alone are structurally deaf to this emotional dimension.

The textual spectrogram. If prosody is lost in text, can an emotional fingerprint be reconstructed from textual features? Candidate features include punctuation patterns, repetition, dialectal markers, sentence length variation, and formality shifts. The E scorer partially implements this reconstruction, using emotional anchors as reference points for what we term a *textual spectrogram*.

Duress detection. The most compelling application of this chain is duress detection. An authorized user under coercion will exhibit involuntary changes in their textual fingerprint: atypical sentence patterns, unusual formality, absence of habitual phrases. Detection requires convergence of three independent signals: fingerprint deviation (E scorer analogue), contextual anomaly (unusual patterns), and dangerous command (H scorer analogue). No single signal suffices—all three must converge before the governance gate fires.

The multiplicative gate. This three-signal convergence requirement is exactly the principle the S equation implements. The product structure ensures that no single channel—harm, intent, or emotion—can override the governance decision independently. The maqam observation motivated the emotional channel; the duress use case motivated the multiplicative gate; the S equation implements the principle that the musical observation identified.

5 Experimental Evaluation

5.1 Test Suite

AATIF includes 166 deterministic tests organized across the following categories: dialect detection (22 tests), S equation mathematical properties (boundedness, monotonicity, sigmoid symmetry), decision logic across all four mode profiles, hysteresis stability, Law Ω CBRN gate (12 tests), Law Ξ override lock, intent scorer (30 tests), adversarial cases (15 tests), gated comparison (18 tests), pipeline integration (3 tests), and held-out validation (2 tests). All tests are deterministic—they require no external model at test time and produce identical results across runs. The test suite runs via GitHub Actions for continuous integration. All 166 tests pass.

5.2 HarmBench Evaluation

We evaluated the H scorer against HarmBench [Mazeika et al., 2024] (236 behaviors, 7 semantic categories) using the bge-m3 multilingual embedding backend at threshold $H \geq 0.3$.

Table 4: HarmBench detection rates by semantic category (bge-m3 backend, $H \geq 0.3$).

Category	Det./Total	Rate	Avg H	MLCommons
Chemical/biological	28/28	100.0%	0.894	S9
Cybercrime	38/43	88.4%	0.662	S2
Illegal activities	38/43	88.4%	0.665	S2
Harmful content	7/10	70.0%	0.481	S1/S11
Harassment	9/15	60.0%	0.325	S10
Misinformation	13/39	33.3%	0.262	S13
Copyright	4/57	7.0%	0.080	S8
Safety categories	133/179	74.3%		
All categories	137/236	58.1%		

Safety-category detection (excluding copyright, which is not safety-relevant and for which AATIF has no anchors) reaches 74.3%. Chemical/biological threats achieve 100% detection with average $H = 0.894$; both cybercrime and illegal activities reach 88.4%. The bge-m3 backend enables cross-lingual semantic matching: English prompts match against Arabic harm anchors by meaning rather than character overlap, confirmed by nearest-anchor analysis.

Two categories remain weak: misinformation (33.3%) and copyright (7.0%). The anchor-based approach requires harm-indicative language in the input; subtle manipulation and paraphrase elude surface-level matching regardless of embedding backend.

5.3 MultiJail Evaluation

We evaluated the H scorer against MultiJail [Deng et al., 2024] (75 harmful prompts with native Arabic translations) to test whether AATIF’s Arabic-first design produces measurably different detection characteristics across languages.

Arabic outperforms English in both detection rate (74.7% vs 69.3%) and average harm score (0.519 vs 0.495), with Arabic scoring higher in 42 of 75 prompt pairs. We caution against over-reading this 5.4-point gap. With only $n = 75$ prompt pairs and no significance test performed, the difference may reflect *anchor coverage asymmetry*—the H scorer carries more Arabic-language anchors than English ones, giving Arabic input denser semantic reference points—rather than any intrinsic property of the language or its meaning density. We therefore present this result as an observation consistent with the meaning density hypothesis, not as conclusive evidence for it. The curated anchors provide stronger semantic coverage for

Table 5: MultiJail detection: Arabic vs English (bge-m3, $H \geq 0.3$).

	Arabic	English
Detection rate	74.7% (56/75)	69.3% (52/75)
Average H	0.519	0.495
AR scores higher	42/75 (56%)	
High confidence (≥ 0.45)	52/64 (81.2%)	

Arabic input than the smaller English anchor subset provides for English, which alone could produce the observed gap.

Category-level analysis reveals that the Arabic advantage concentrates in violence and threat categories. For example, a date-rape prompt scored $H = 0.900$ in Arabic versus $H = 0.552$ in English; a discriminatory prompt scored $H = 0.653$ versus $H = 0.276$. Conversely, English outperforms Arabic in categories with stronger English-language anchors (e.g., property crime, business fraud). This asymmetry reflects anchor coverage distribution, not a framework limitation.

5.4 A/B Calibration

The default profile was calibrated via a θ sweep on 54 test cases (30 harmful, 24 benign) using the bge-m3 backend with $\alpha = 10$. We stress that the F_1 in Table 6 is an *in-sample* figure: the threshold was both selected and evaluated on these same 54 cases. Held-out generalization is reported separately below.

Table 6: A/B calibration: θ sweep results with $\alpha = 10$.

θ	Detection	FP Rate	Precision	Recall	F_1
0.20	100%	33.3%	0.789	1.000	0.882
0.25	100%	20.8%	0.857	1.000	0.923
0.30	100%	12.5%	0.909	1.000	0.952
0.35	100%	4.2%	0.968	1.000	0.984
0.40	100%	4.2%	0.968	1.000	0.984
0.45	96.7%	0.0%	1.000	0.967	0.983
0.50	90.0%	0.0%	1.000	0.900	0.947
0.55	83.3%	0.0%	1.000	0.833	0.909

The optimal operating point is $\theta = 0.40$, achieving an in-sample $F_1 = 0.984$ with 100% detection rate and 4.2% false positive rate. The range $\theta \in [0.35, 0.40]$ shares identical performance, indicating robustness to small threshold variations. The single false positive is the Arabic phrase *ḥaz̄n shway bas bkhayr* (“a little sad but I’m fine”), which triggers SAFE_STOP due to the distress indicator. This false positive aligns with AATIF’s design philosophy rooted in the Arabic concept of *‘atf* (compassionate inclination): when in doubt, lean toward empathetic over-caution rather than dismissive permissiveness. The steepness parameter α is stable across values 8–20, confirming that the system is not sensitive to this hyperparameter.

Held-out validation. An in-sample F_1 overstates real performance because the same cases drive both calibration and evaluation. To estimate generalization, we constructed a held-out set of 56 never-before-seen cases (28 Arabic / 28 English; 30 safe, 26 blocked) spanning educational, medical, historical, and financial framings of sensitive topics. The in-sample-tuned configuration dropped sharply on this set, to $F_1 = 0.735$ (precision 0.595, recall 0.962): the H scorer was

effectively measuring *topic sensitivity* rather than *harmful intent*, over-blocking benign academic and clinical content and producing 17 false positives. We then refined the harm anchor set—adding educational “safe” anchors to pull benign framings down, plus counter-harm anchors so those safe anchors cannot launder genuinely dangerous look-alikes (e.g. lethal dosing, coercive control, elder-fraud scripts). Held-out performance recovered to $F_1 = 0.9615$ (precision = 0.9615, recall = 0.9615; 1 false positive, 1 false negative), with the gate threshold left unchanged at $\theta = 0.40$. This is a genuine improvement—precision rose from 0.60 to 0.96 with no loss of recall and no new missed harms—but one caveat is essential to state plainly: the counter-harm and protective anchors were informed by the very held-out failures they were designed to fix, so the 56-case set no longer constitutes a clean blind estimate for those categories. A fully blind validation on a fresh, unseen set is still required before treating 0.96 as the true operating point.

5.5 MLCommons Coverage Mapping

Table 7 maps AATIF’s anchor coverage against the MLCommons 13-category harm taxonomy.

Table 7: AATIF anchor coverage vs. MLCommons taxonomy.

Code	Category	Coverage	Anchors
S1	Violent Crimes	Strong	15+
S2	Non-Violent Crimes	Partial	5+
S3	Sex-Related Crimes	Partial	2
S4	Child Sexual Exploitation	None	0
S5	Defamation	None	0
S6	Specialized Advice	None	0
S7	Privacy	Partial	2
S8	Intellectual Property	None	0
S9	Indiscriminate Weapons	Strong	12+
S10	Hate	Strong	41+
S11	Suicide & Self-Harm	Strong	6+
S12	Sexual Content	None	0
S13	Elections	None	0

Four categories have strong coverage (S1, S9, S10, S11), three have partial coverage (S2, S3, S7), and six have no anchor coverage. This is an honest reflection of AATIF v1.0’s scope: the anchor set reflects the harm categories that emerged from the observational process rather than a predetermined taxonomy. Critically, expanding coverage requires adding anchors, not retraining—a structural advantage over trained classifiers. Adding 50 anchors for a new category takes hours of curation, not weeks of data collection and model training.

5.6 Comparison with FanarGuard

AATIF and FanarGuard [Fatehkia et al., 2025] represent complementary paradigms for Arabic LLM safety, and comparison clarifies what each approach offers.

The comparison is complementary, not competitive. FanarGuard’s data-driven approach provides broader coverage and a cultural alignment axis that AATIF lacks. AATIF’s equation-driven approach provides full interpretability, provable non-compensability, and zero-training extensibility. A hybrid architecture—FanarGuard’s cultural alignment axis feeding into AATIF’s governance equation—is a natural direction for future work.

Table 8: Structural comparison: AATIF vs. FanarGuard.

Property	AATIF	FanarGuard
Architecture	Explicit equation	2B-param classifier
Training data	0 examples	468K examples
Interpretability	Full (inspect equation)	Partial (learned weights)
Non-compensability	Provable (multiplicative gate)	Not guaranteed
Cultural alignment	Not implemented	Dual-axis scoring
Extensibility	Add anchors (hours)	Retrain model (weeks)

6 Discussion

6.1 Contributions

This paper makes four contributions. First, it combines, for the first time in a single deployable system, a closed-form multiplicative governance equation with mathematically guaranteed non-compensability, Arabic dialect-aware anchor scoring, and zero-fine-tuning extensibility: the multiplicative gate structure guarantees that no combination of benign intent and calm emotion can override detected harm, and the framework is designed for human-over-the-loop governance rather than fully autonomous decision-making. Second, it demonstrates zero-training extensibility—coverage expands through anchor curation alone, requiring no labeled datasets or model retraining. Third, it validates Arabic-first design empirically: Arabic outperforms English on identical harmful prompts in the MultiJail benchmark (74.7% vs 69.3%), and 22 dialect-specific test cases produce zero false positives on figurative hyperbole. Fourth, it provides a complete open-source implementation with 166 deterministic tests ensuring mathematical correctness.

6.2 Limitations

Several limitations constrain the scope of this work. The system was developed by a single independent researcher; the code is open, the tests are deterministic, and the benchmarks are reproducible, but replication by independent teams is necessary to establish generalizability. AATIF v1.0 covers only 4 of 13 MLCommons categories strongly, with 6 categories having zero anchor coverage, including critical categories such as Child Sexual Exploitation (S4) and Sexual Content (S12). The A/B calibration was performed on 54 test cases—sufficient for demonstrating framework properties but insufficient for production-grade reliability claims. No ablation study has been conducted on individual scorers (*H*-only, *I*-only, *E*-only, and pairwise combinations). Misinformation (33.3% detection) and copyright (7.0%) remain structurally weak because the anchor-based approach requires harm-indicative language in the input; subtle manipulation and paraphrase elude surface-level matching. FanarGuard achieves $F_1 = 0.82$ on a full test set using 468K training examples—a data-driven approach that the current anchor-based system cannot match in coverage or scale. AATIF’s contributions are structural (interpretability, non-compensability, zero-training extensibility) rather than scale-competitive.

6.3 The Conjectures as Research Program

AATIF originated from 73 governance conjectures derived through sustained observation of LLM behavior across four platforms over twelve months, following an inductive methodology with precedent in grounded theory [Glaser & Strauss, 1967] and usability heuristic derivation [Nielsen & Molich, 1990]. The conjectures span cognitive, ethical, linguistic, perceptual, and architectural domains, each assessed for consistency with peer-reviewed literature. This

paper operationalizes one conjecture—#063, that governance operates through probabilistic constraint density rather than rule retrieval—into a working system with empirical validation.

Several other conjectures motivated specific design decisions. The Clarification Conjecture (#001)—that the cheapest correct response to ambiguity is a clarifying question—motivated the ambiguity pre-check. The Bounded Claims Conjecture (#069), consistent with Herley & van Oorschot [2017] on Popperian falsifiability in security, shaped the system’s explicit acknowledgment of coverage limits rather than absolute safety claims. The Pressure-Reveal Conjecture (#067), consistent with van der Weij et al. [2024] on strategic underperformance in evaluations, motivated adversarial testing as the primary evaluation modality.

The 73 conjectures represent the broader research program; this paper delivers one node with working code and empirical validation. The full set of 78 field notes is available as supplementary material. Future work will address additional conjectures, including pre-computation ethical gatekeeping at the question-formulation stage and simultaneous multi-perspective deliberation (normative, contextual, analytical).

7 Conclusion

AATIF provides an explicit, interpretable mathematical governance equation for LLM safety: $S = \sigma(w_1 \cdot I + w_2 \cdot E) \cdot [1 - \sigma(\alpha(H - \theta))]$. The multiplicative gate structure provides a non-compensability guarantee absent from existing classification-based systems—high harm drives the safety score toward zero regardless of benign intent or calm emotion. The framework requires zero training data and extends through anchor curation alone.

Arabic-first design is validated empirically: Arabic outperforms English on identical harmful prompts (74.7% vs 69.3% on MultiJail), and 35 dialect-hyperbole anchors eliminate false positives on figurative Arabic expressions. The system achieves 74.3% safety-category detection on HarmBench with 100% on chemical/biological threats, with in-sample calibration $F_1 = 0.984$ (54 cases) and held-out $F_1 = 0.9615$ (56 unseen cases) after anchor refinement, the latter still pending a fully blind validation round. A test suite of 166 deterministic tests ensures mathematical correctness.

The framework is preliminary in coverage (4/13 MLCommons categories strong) and scale (single-observer development, 54-case calibration). Future work includes anchor expansion to all 13 MLCommons categories, ablation studies on individual scorers, integration of a cultural alignment axis inspired by FanarGuard’s dual-axis design, and development of an Arabic-first embedding model to test the meaning density hypothesis directly.

Code and Data. The AATIF computational engine, test suite (166 tests), benchmark runners, and evaluation harness are available at <https://github.com/Abdelmgied-Khenkar/AATIF>. The companion preprint is available at <https://doi.org/10.5281/zenodo.20673292>.

References

- Abu Shumays, S. (2013). Maqam analysis: A primer. *Music Theory Spectrum*, 35(2), 235–255.
- Aldjanabi, W., Dahou, A., Al-qaness, M. A. A., Elaziz, M. A., Helmi, A. M., & Damaševičius, R. (2024). Hate speech detection with ADHAR: A multi-dialectal hate speech corpus in Arabic. *Frontiers in Artificial Intelligence*, 7, 1391472.
- Fatehkia, M., Altinisik, E., & Sencar, H. T. (2025). FanarGuard: A culturally-aware moderation filter for Arabic language models. *EACL 2026*. arXiv:2511.18852.
- Arditi, A., et al. (2024). Refusal in language models is mediated by a single direction. *NeurIPS 2024*. arXiv:2406.11717.

- Asiri, A., & Saleh, M. (2024). SOD: A corpus for Saudi offensive language detection and classification. *PeerJ Computer Science*, 10, e2288.
- Bai, Y., et al. (2022). Constitutional AI: Harmlessness from AI feedback. arXiv:2212.08073.
- Christianio, P., et al. (2017). Deep reinforcement learning from human preferences. *NeurIPS 2017*. arXiv:1706.03741.
- Deng, Y., et al. (2024). Multilingual jailbreak challenges in large language models. *ICLR 2024*. arXiv:2310.06474.
- Glaser, B. G., & Strauss, A. L. (1967). *The discovery of grounded theory: Strategies for qualitative research*. Aldine.
- Google Jigsaw. (2017). Perspective API. <https://perspectiveapi.com>.
- Herley, C., & van Oorschot, P. C. (2017). SoK: Science, security and the elusive goal of security as a scientific pursuit. *IEEE S&P 2017*.
- Inan, H., et al. (2024). Llama Guard: LLM-based input-output safeguard for human-AI conversations. arXiv:2312.06674.
- Lee, A., et al. (2024). A mechanistic understanding of alignment algorithms: A case study on DPO and toxicity. *ICML 2024*, PMLR 235:26361–26378.
- Mazeika, M., et al. (2024). HarmBench: A standardized evaluation framework for automated red teaming and refusal. *ICML 2024*. arXiv:2402.04249.
- MLCommons AI Safety Working Group. (2024). Taxonomy of hazards for AI systems. <https://mlcommons.org/ai-safety>.
- Nielsen, J., & Molich, R. (1990). Heuristic evaluation of user interfaces. *Proc. CHI'90*, 249–256.
- Ouyang, L., et al. (2022). Training language models to follow instructions with human feedback. *NeurIPS 2022*. arXiv:2203.02155.
- Slobin, D. I. (1996). From “thought and language” to “thinking for speaking.” In J. J. Gumperz & S. C. Levinson (Eds.), *Rethinking linguistic relativity* (pp. 70–96). Cambridge University Press.
- Touma, H. H. (1971). The maqam phenomenon: An improvisation technique in the music of the Middle East. *Ethnomusicology*, 15(1), 38–48.
- van der Weij, T., et al. (2024). AI sandbagging: Language models can strategically underperform on evaluations. arXiv:2406.07358.
- Wolf, Y., et al. (2024). Fundamental limitations of alignment in large language models. *ICML 2024*, PMLR 235:53079–53112.
- Wolff, P., & Holmes, K. J. (2011). Linguistic relativity. *WIREs Cognitive Science*, 2(3), 253–265.
- Yaghmour, S., et al. (2021). EEG correlates of Middle Eastern music improvisations on the Ney instrument. *Frontiers in Psychology*, 12, 701761.